

In [20]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

In [21]:

```
df=pd.read_csv(r"C:\Users\RAMADEVI SURIPAKA\Downloads\bottle.csv.zip")
df
```

C:\Users\RAMADEVI SURIPAKA\AppData\Local\Temp\ipykernel_5452\1481418663.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.
df=pd.read_csv(r"C:\Users\RAMADEVI SURIPAKA\Downloads\bottle.csv.zip")

Out[21]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	R_PHAEO	R_PRES	R_SAMP	
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN	25.64900	NaN	...	NaN	0	NaN	
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN	25.65600	NaN	...	NaN	8	NaN	
2	1	3	054.0 056.0	19-4903CR-HY-060-0930-05400560-0010A-7	10	10.460	33.4370	NaN	25.65400	NaN	...	NaN	10	NaN	
3	1	4	054.0 056.0	19-4903CR-HY-060-0930-05400560-0019A-3	19	10.450	33.4200	NaN	25.64300	NaN	...	NaN	19	NaN	
4	1	5	054.0 056.0	19-4903CR-HY-060-0930-05400560-0020A-7	20	10.450	33.4210	NaN	25.64300	NaN	...	NaN	20	NaN	
...
864858	34404	864859	093.4 026.4	20-1611SR-MX-310-2239-09340264-0000A-7	0	18.744	33.4083	5.805	23.87055	108.74	...	0.18	0	NaN	
864859	34404	864860	093.4 026.4	20-1611SR-MX-310-2239-09340264-0002A-3	2	18.744	33.4083	5.805	23.87072	108.74	...	0.18	2	4.0	
864860	34404	864861	093.4 026.4	20-1611SR-MX-310-2239-09340264-0005A-3	5	18.692	33.4150	5.796	23.88911	108.46	...	0.18	5	3.0	
864861	34404	864862	093.4 026.4	20-1611SR-MX-310-2239-09340264-0010A-3	10	18.161	33.4062	5.816	24.01426	107.74	...	0.31	10	2.0	
864862	34404	864863	093.4 026.4	20-1611SR-MX-310-2239-09340264-0015A-3	15	17.533	33.3880	5.774	24.15297	105.66	...	0.61	15	1.0	

864863 rows × 74 columns

In [22]:

```
df=df[['Salnty','T_degC']]
df.columns=['Sal','Temp']
```

In [23]:

```
df.head()
```

Out[23]:


	Sal	Temp
0	33.440	10.50
1	33.440	10.46
2	33.437	10.46
3	33.420	10.45
4	33.421	10.45

In [24]:

```
sns.lmplot(x='Sal',y='Temp',data=df,order=2,ci=None)
```

Out[24]:

<seaborn.axisgrid.FacetGrid at 0x1d39fccfaf0>



In [25]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  --
 0    Sal      817509 non-null  float64
 1    Temp     853900 non-null  float64
dtypes: float64(2)
memory usage: 13.2 MB
```

In [26]:

df.describe()

Out[26]:

	Sal	Temp
count	817509.000000	853900.000000
mean	33.840350	10.799677
std	0.461843	4.243825
min	28.431000	1.440000
25%	33.488000	7.680000
50%	33.863000	10.060000
75%	34.196900	13.880000
max	37.034000	31.140000

In [27]:

df.fillna(method='ffill')

Out[27]:

	Sal	Temp
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

In [31]:

df.isnull().sum()

Out[31]:

```
Sal      0
Temp     0
dtype: int64
```

In [28]:

df.fillna(value=0,inplace=True)

C:\Users\RAMADEVI SURIPAKA\AppData\Local\Temp\ipykernel_5452\1434098079.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.fillna(value=0,inplace=True)

In [32]:

```
x=np.array(df['Sal']).reshape(-1,1)
y=np.array(df['Temp']).reshape(-1,1)
```

In [33]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
reg=LinearRegression()
reg.fit(x_train,y_train)
print(reg.score(x_test,y_test))
```

0.014132816960575889

In [34]:

```
df.isna().any()
```

Out[34]:

```
Sal      False
Temp     False
dtype: bool
```

In [35]:

```
df.isnull().sum()
```


Out[35]:

```
Sal      0
Temp     0
dtype: int64
```

In [36]:

```
y_pred=reg.predict(x_test)

plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```




In [37]:

```
df500=df[:][:500]
sns.lmplot(x='Sal',y='Temp',data=df500,order=1,ci=None)
```

Out[37]:


<seaborn.axisgrid.FacetGrid at 0x1d39fd4fbe0>



In [38]:

```
df500.fillna(method='ffill',inplace=True)
x=np.array(df500['Sal']).reshape(-1,1)
y=np.array(df500['Temp']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
reg=LinearRegression()
reg.fit(x_train,y_train)
print("Regresion:",reg.score(x_test,y_test))
y_pred=reg.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regresion: 0.04946978366895738



In [39]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

In [40]:

```
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.04946978366895738

IMPLEMENTATION RIDGE AND LASSO

In [41]:


```
from sklearn.linear_model import Ridge
from sklearn.linear_model import RidgeCV
from sklearn.linear_model import Lasso
```

In [42]:

```
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True)
```

Out[42]:

<Axes: >



In [43]:

```
features = df.columns[0:2]
target = df.columns[-1]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X_train is (605404, 2)
 The dimension of X_test is (259459, 2)

In [44]:

```
lr = LinearRegression()  
#Fit model  
lr.fit(X_train, y_train)  
#predict  
#prediction = lr.predict(X_test)  
#actual  
actual = y_test  
train_score_lr = lr.score(X_train, y_train)  
test_score_lr = lr.score(X_test, y_test)  
print("\nLinear Regression Model:\n")  
print("The train score for lr model is {}".format(train_score_lr))  
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0
The test score for lr model is 1.0

In [45]:

```
ridgeReg=Ridge(alpha=10)  
ridgeReg.fit(X_train,y_train)  
train_score_ridge=ridgeReg.score(X_train,y_train)  
test_score_ridge=ridgeReg.score(X_test,y_test)  
print("\nRidge Model:\n")  
print("The train score for ridge model is {}".format(train_score_ridge))  
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.999999999723243
The test score for ridge model is 0.9999999997231402

In [49]:

```
.figure(figsize=(10,10))

.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha=10$',zorder=1)
t.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge;$\alpha=100$',zorder=2)
.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
.xticks(rotation=90)
.legend()
t.title("comparison plot of Ridge,Lasso and Linear regression model")
.show()
```



In [47]:

```
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:


The train score for ls model is 0.0
 The test score for ls model is -1.9031696447013857e-05

In [48]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[48]:

<Axes: >



In [51]:

```
from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train, y_train)
print(lasso_cv.score(X_train,y_train))
print(lasso_cv.score(X_test,y_test))
```

0.9999999994806811


0.9999999994806712

In [54]:

```

t.figure(figsize=(10,10))
t.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;\alpha=10$',zorder=1)
t.plot(lasso_cv.coef_,alpha=0.6,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge;\alpha=grid$')
t.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
t.xticks(rotation=90)
t.legend()
t.title("comparison plot of Ridge,Lasso and Linear regression model")
t.show()

```



ELASTICNET

In [55]:

```

from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)

```

```

[0.          0.94934511]
0.540121963106797

```

In [56]:

```
y_pred_elastic=regr.predict(X_train)
```

In [57]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)  
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 114.40984808659212

In []: